


9. Hashing

- ◆ Hashing as a method of indexing.
- ◆ The components of a hash table.
- ◆ Hash function, how it works, design of a hash function.
- ◆ Collision, what it is, the techniques for handling collisions, their advantages and disadvantages.

Hash tables

1

Hash Tables



Hashing:

- ◆ Hashing is a method for **directly** referencing items in a dictionary (data repository) by doing **arithmetic transformations** on **keys** into dictionary **addresses**.
- ◆ A **hash function** is perfect if there is no key **collision**, that is, two keys hash to the same hash value (address).

Hash tables

3

Why Hash Tables?


- ◆ All search structures so far
 - Relied on a comparison operation
 - Performance $O(n)$ or $O(\log n)$
- ◆ Assume we have a function
 - $f(\text{key}) \rightarrow \text{integer}$
ie one that maps a key to an integer
- ◆ What performance might we expect now?

Hash tables

4

Hash Tables - Structure

- ◆ Simplest case:
 - Assume items have integer keys in the range $1 \dots m$
 - Use the value of the key itself to select a slot in a **direct access table** to **store** the item
 - To **search** for an item with key, k , just look in slot k
 - If there's an item there, you've found it
 - If the tag is 0, it's missing.
 - Constant time, $O(1)$



Hash tables

5

Hash Tables - Constraints


- ◆ Constraints
 - Keys must be unique
 - Keys must be integers
 - Keys must lie in a small range
 - For storage efficiency, keys must be **dense** in the range
 - If they're **sparse** (lots of gaps between values), a lot of (unnecessary) space is used to obtain speed
 - Space for speed trade-off

Hash tables

6

Hash Tables - Relaxing the constraints

- ◆ Keys must be unique
 - Construct a linked list of duplicates : "attached" to each slot
 - If a search can be satisfied by **any** item with key, k , performance is still $O(1)$ *but*
 - If the item has some other distinguishing feature which must be matched, we get $O(n_{max})$, where n_{max} is the largest number of duplicates - or length of the longest chain




Hash tables

7

Hash Tables - Relaxing the constraints

- ◆ Keys are integers
 - Need a **hash function** $h(\text{key}) \rightarrow \text{integer}$ *ie* one that maps a key of a different type (e.g. char) to an integer
 - Applying this function to the key produces an address
 - If h maps each key to a **unique integer** in the range $0 \dots m-1$, then search is $O(1)$



Hash tables

8

An Example: Perfect Hash

- ◆ suppose: MagicNumber = 15
 - ◆

```
int h(String s) {
    return ((s[0] + s[1])% MagicNumber);
}
```
 - ◆ suppose:


```
typedef struct {
    String name;
    int numMoons;
    double sunDistance;
} planet;
```
- planet solarSystem[MagicNumber];

Hash tables

9

Suppose:

```
solarSystem[h("Mercury")] = {"Mercury", 0, 36.0};
solarSystem[h("Venus")] = {"Venus", 0, 67.27};
solarSystem[h("Earth")] = {"Earth", 1, 93.0};
solarSystem[h("Mars")] = {"Mars", 2, 141.71};
solarSystem[h("Jupiter")] = {"Jupiter", 16, 483.88};
solarSystem[h("Saturn")] = {"Saturn", 12, 887.14};
solarSystem[h("Uranus")] = {"Uranus", 5, 1783.98};
solarSystem[h("Neptune")] = {"Neptune", 2, 2795};
solarSystem[h("Pluto")] = {"Pluto", 1, 3675};
```

Where are they located

Saturn	Earth	Uranus	Venus	Pluto	Mars	Jupiter	Mercury	Neptune
0	1	2	3	4	5	6	7	8
9	10	11	12	13	14			

Hash tables

10

"Ju" in ASCII are 74 and 117, $74 + 117 = 191$;
 $191 \% 15 = 11$;

```
h("Mercury") = 13
h("Venus") = 7
h("Earth") = 1
h("Mars") = 9
h("Jupiter") = 11
h("Saturn") = 0
h("Uranus") = 4
h("Neptune") = 14
h("Pluto") = 8
```

Thus, our search function is simply:
 planet search(String s){ return solarSystem[h(s)]; }

Hash tables

11

Hash Functions

- ◆ A **hash function** h maps keys of a given type to integers in a fixed interval $[0, N - 1]$
- ◆ Example:

$$h(x) = x \bmod N$$
 is a hash function for integer keys
- ◆ The integer $h(x)$ is called the **hash value** of key x
- ◆ The goal of a hash function is to uniformly disperse keys in the range $[0, N - 1]$

Hash tables

12

Choosing the Hash Function

Uniform hashing

Ideal hash function

- $P(k)$ = probability that a key, k , occurs
- If there are m slots in our hash table,
- a **uniform hashing function**, $h(k)$, would ensure:

$$\sum_{k/h(k)=0} P(k) = \sum_{k/h(k)=1} P(k) = \dots = \sum_{k/h(k)=m-1} P(k) = \frac{1}{m}$$

Read as sum over all k such that $h(k) = 0$

- or, in plain English,
- the number of keys that map to each slot is equal

Hash tables

13

Hash Tables - A Uniform Hash Function

- ◆ If the keys are integers randomly distributed in $[0, r)$, then

$$h(k) = \left\lfloor \frac{mk}{r} \right\rfloor$$

is a **uniform hash function**

- ◆ Most hashing functions can be made to map the keys to $[0, r)$ for some r
 - e.g. adding the ASCII codes for characters mod 255 will give values in $[0, 256)$ or $[0, 255]$

Hash tables

14

Hash Tables


- ◆ A **hash table** for a given key type consists of
 - Hash function h
 - Array (called table) of size N
- ◆ When implementing a dictionary with a hash table, the goal is to store item (k, o) at index $i = h(k)$
- ◆ A **collision** occurs when two keys in the dictionary have the same hash value, i.e., $h(k) = h(k')$, whereas $k \neq k'$
- ◆ Collision handling schemes:
 - **Chaining**: colliding items are stored in a sequence
 - **Open addressing**: the colliding item is placed in a different cell of the table

Hash tables

15

Example

- ◆ We design a hash table for a dictionary storing items (Phone#, Name), where a Phone# is a ten-digit positive integer
- ◆ Our hash table uses an array of size $N = 10,000$ and the hash function $h(x) = \text{last four digits of } x$
- ◆ We use chaining to handle collisions



Hash tables

16

Define Hash Functions

- ◆ A hash function is usually specified as the composition of two functions:
 - Hash code mapping:** h_1 : keys \rightarrow integers
 - Compression mapping:** h_2 : integers $\rightarrow [0, N - 1]$
- ◆ The hash code mapping is applied first, and the compression mapping is applied next on the result, i.e., $h(x) = h_2(h_1(x))$
- ◆ The goal of the hash function is to "disperse" the keys in an apparently random way

Hash tables

17

Hash Code Mappings

- ◆ **Integer cast:**
 - We reinterpret the bits of the key as an integer
 - Suitable for keys of length less than or equal to the number of bits of the integer type (e.g., byte, short, int, and long)
- ◆ **Component sum:**
 - Partition the bits of the key into components of fixed length (e.g., 16 or 32 bits) and
 - Sum the components (ignoring overflows)
 - Suitable for numeric keys of fixed length greater than or equal to the number of bits of the integer type (e.g., int and long)

Hash tables

18

Example: A Hash Function

◆ Hash function

- With this hash function

```
int hash( char *s, int n ) {
    int sum = 0;
    while( n-- ) sum = sum + *s++;
    return sum % 256;
}
```

- hash("AB", 2) and hash("BA", 2) return the same value!
- This is called a **collision**
- A variety of techniques are used for resolving collisions

Hash tables

19

Hash Code Mappings (cont.)

◆ Polynomial accumulation:

- We partition the bits of the key into a sequence of components of fixed length (e.g., 8, 16 or 32 bits)

$$a_0 a_1 \dots a_{n-1}$$

- We evaluate the polynomial

$$p(z) = a_0 + a_1 z + a_2 z^2 + \dots + a_{n-1} z^{n-1}$$

at a fixed value z

- Especially suitable for strings (e.g., the choice $z = 33$ gives at most 6 collisions on a set of 50,000 English words)

Hash tables

20

Hash Code Mappings (cont.)

- Polynomial $p(z)$ can be evaluated in $O(n)$ time using Horner's rule:

- The following polynomials are successively computed, each from the previous one in $O(1)$ time

$$p_0(z) = a_{n-1}$$

$$p_i(z) = a_{n-i-1} + z p_{i-1}(z)$$

$$(i = 1, 2, \dots, n-1)$$

- We have $p(z) = p_{n-1}(z)$

```
int poly(int a[], int z; int n){
    int p = 0;
    for (int i = n-1; i >= 0; i--){
        p = a[i] + z*p;
    }
    return p;
}
```

Hash tables

21

Compression Mappings

◆ Division:

- $h_2(y) = y \bmod N$
- The size N of the hash table is usually chosen to be a prime
- The reason has to do with number theory and is beyond the scope of this course

◆ Multiply, Add and Divide (MAD):

- $h_2(y) = (ay + b) \bmod N$
- a and b are nonnegative integers such that $a \bmod N \neq 0$
- Otherwise, every integer would map to the same value b

Hash tables

22

Linear Probing for handling collision

- Linear probing is a method of open addressing

- Linear probing handles collisions by placing the colliding item in the next (circularly) available table cell

- Each table cell inspected is referred to as a "probe"

- Colliding items lump together. Future collisions may cause a longer sequence of probes

◆ Example:

- $h(x) = x \bmod 13$
- Insert keys 18(5), 41(2), 22(9), 44(5), 59(7), 32(6), 31(5), 73(8), in this order

0	1	2	3	4	5	6	7	8	9	10	11	12
		41		18	44	59	32	22	31	73		
0	1	2	3	4	5	6	7	8	9	10	11	12

Hash tables

23

Search with Linear Probing

- Consider a hash table A that uses linear probing

◆ findElement(k)

- We start at cell $h(k)$
- We probe consecutive locations until one of the following occurs
 - An item with key k is found, or
 - An empty cell is found, or
 - N cells have been unsuccessfully probed

```
function findElement(k){
    i = h(k);
    p = 0;
    repeat {
        c = A[i];
        if (c == 0)
            return NO_SUCH_KEY;
        else if (c.key == k)
            return c.element();
        else {
            i = (i + 1) mod N;
            p = p + 1;
        }
    } until (p == N);
    return NO_SUCH_KEY;
}
```

Hash tables

24

Updates with Linear Probing

- ◆ To handle insertions and deletions, we introduce a special key flag, called *AVAILABLE*, which replaces deleted elements
- ◆ **removeElement(*k*)**
 - We search for an item with key *k*
 - If such an item (*k*, *o*) is found, we replace it with the special item *AVAILABLE* and we return element *o*
 - Else, we return *NO_SUCH_KEY*
- ◆ **insert Item(*k*, *o*)**
 - We report an error if the table is full
 - We start at cell *h(k)*
 - We probe consecutive cells until one of the following occurs
 - A cell *i* is found that is either empty or stores *AVAILABLE*, or
 - *N* cells have been unsuccessfully probed
 - We store item (*k*, *o*) in cell *i*

Hash tables

25

Double Hashing for handling collision

- ◆ Double hashing uses a secondary hash function *d(k)* and handles collisions by placing an item in the first available cell of the series $(i + jd(k)) \bmod N$ for $j = 0, 1, \dots, N-1$
- ◆ The secondary hash function *d(k)* cannot have zero values
- ◆ The table size *N* must be a prime to allow probing of all the cells
- ◆ Common choice of compression map for the secondary hash function: $d_2(k) = q - k \bmod q$ where
 - $q < N$
 - *q* is a prime
- ◆ The possible values for *d₂(k)* are 1, 2, ..., *q*

Hash tables

26

Example of Double Hashing

- ◆ Consider a hash table storing integer keys that handles collision with double hashing
 - $N = 13$
 - $h(k) = k \bmod 13$
 - $d(k) = 7 - k \bmod 7$
- ◆ Insert keys 18, 41, 22, 44, 59, 32, 31, 73, in this order

<i>k</i>	<i>h(k)</i>	<i>d(k)</i>	Probes
18	5	3	5
41	2	1	2
22	9	6	9
44	5	5	5 10
59	7	4	7
32	6	3	6
31	5	4	5 9 0
73	8	4	8

0	1	2	3	4	5	6	7	8	9	10	11	12

↓

31	41				18	32	59	73	22	44		
0	1	2	3	4	5	6	7	8	9	10	11	12

Hash tables

27

Performance of Probing:

- ◆ Let *N* be the number of slots of a hash table, *n* be the number of items in the table, we define load factor as:

$$\alpha = n/N$$
- ◆ If the hash function randomly distributes keys through the table, then the expected length of a successful search path is:

$$\text{length}_{\text{succ}} = \frac{1}{2} (1 + 1/(1 - \alpha))$$

Hash tables

28

Performance of Probing:

- ◆ The expected length of an unsuccessful search is approximately:

$$\text{length}_{\text{unsucc}} = \frac{1}{2} (1 + 1/(1 - \alpha)^2)$$

Hash tables

29

Problems with Probing:


- ◆ The size of the hash table must be fixed in advance.
- ◆ The search costs increase dramatically as the table becomes nearly full.
- ◆ Need a special object, called *AVAILABLE*, to implement "delete" operation.

Hash tables

30

Collision resolution using linked Lists:

- ◆ Dynamically allocate space.
- ◆ Easy to insert/delete an item
- ◆ Need a link for each node in the hash table.
- ◆ Also called chaining.



Hash tables

31

Performance:

- ◆ Let N be the size of the hash table, n the number of items in the table's linked lists, if all input sequences are equally likely and the hash function randomly distributes keys over the table, the expected length of a linked list is n/N .

$$\text{length}_{\text{succ}} = 1 + \alpha/2$$


$$\text{length}_{\text{unsucc}} = \alpha$$

Hash tables

32

Collision resolution using overflow area

- ✕ Overflow area
 - Linked list constructed in special area of table called **overflow area**
- $h(k) == h(j)$
- k stored first
- Adding j
 - Calculate $h(j)$
 - Find k
 - Get first slot in overflow area
 - Put j in the slot
 - k 's pointer points to this slot
- Searching - same as linked list



Hash tables

33

Collision Resolution Summary

- ◆ Probing
 - Fixed number of elements
 - Multiple collisions become probable
 - Search costs increase dramatically as the table becomes nearly full.
- ◆ Chaining
 - Unlimited number of elements. Unlimited number of collisions
 - Overhead of multiple linked lists
- ◆ Re-hashing
 - Maximum number of elements must be known
 - Multiple collisions become probable
- ◆ Overflow area
 - Collisions don't use primary table space
 - Performance same as linked lists

Hash tables

34

Conclusion:

- ◆ In the worst case, searches, insertions and removals on a hash table take $O(n)$ time
- ◆ The worst case occurs when all the keys inserted into the dictionary collide
- ◆ The load factor $\alpha = n/N$ affects the performance of a hash table
- ◆ Assuming that the hash values are like random numbers, it can be shown that the expected number of probes for an insertion with open addressing is $\frac{1}{2} (1 + 1/(1 - \alpha))$
- ◆ The expected running time of all the dictionary ADT operations in a hash table is $O(1)$
- ◆ In practice, hashing is very fast provided the load factor is not close to 100%
- ◆ Applications of hash tables:
 - small databases
 - compilers
 - browser caches

Hash tables

35

Collision Frequency

- ◆ Birthdays or the von Mises paradox
 - There are 365 days in a normal year
 - ◀ Birthdays on the same day unlikely?
 - How many people do I need before "it's an even bet" (ie the probability is > 50%) that two have the same birthday?
 - View
 - the days of the year as the slots in a hash table
 - the "birthday function" as mapping people to slots
 - Answering von Mises' question answers the question about the probability of collisions in a hash table



Hash tables

36

Distinct Birthdays

- Let $Q(n)$ = probability that n people have **distinct** birthdays

- $Q(1) = 1$

- With two people, the 2nd has only 364 "free" birthdays

$$Q(2) = Q(1) * \frac{364}{365}$$

- The 3rd has only 363, and so on:

$$Q(n) = Q(1) * \frac{364}{365} * \frac{363}{365} * \dots * \frac{365-n+1}{365}$$

Hash tables

37


Coincident Birthdays

- Probability of having two **identical** birthdays

- $P(n) = 1 - Q(n)$

- $P(23) = 0.507$

- With 23 entries, table is only $23/365 = 6.3\%$ full!



Hash tables

38

Hash Tables - Load factor

- Collisions are very probable!

- Table load factor

$$\alpha = \frac{n}{m} \quad \begin{array}{l} n = \text{number of items} \\ m = \text{number of slots} \end{array}$$

must be kept low

- Detailed analyses of the average chain length (or number of comparisons/search) are available

- Separate chaining**

- linked lists attached to each slot gives best performance
- but uses more space!

Hash tables

39