

Process


Process Concepts

- Process Definition
- Process Creation
- **Process States and State Transitions**
- Important Information Associated with Processes



Two-State Process Model

- Process may be in one of two states
 - ⌘ Running
 - ⌘ Not-running

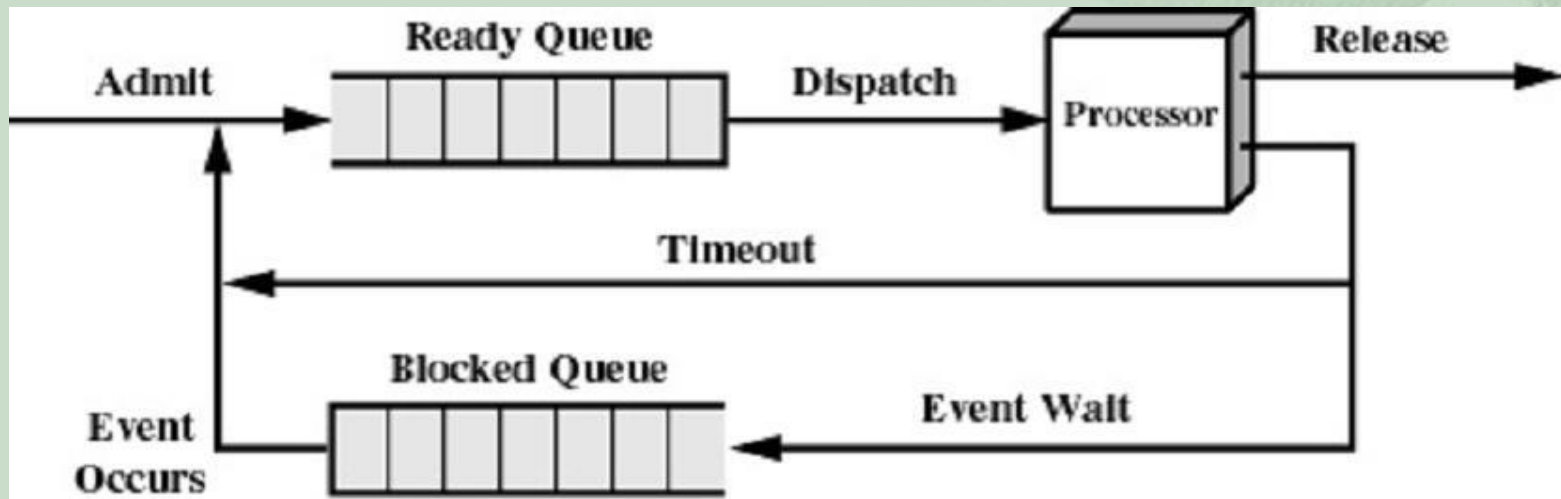



Two-state Model Problems

- Not-running
 - ↳ ready to execute
- Blocked
 - ↳ waiting for I/O, semaphore
- Dispatcher cannot just select the process that has been in the queue the longest because it may be blocked

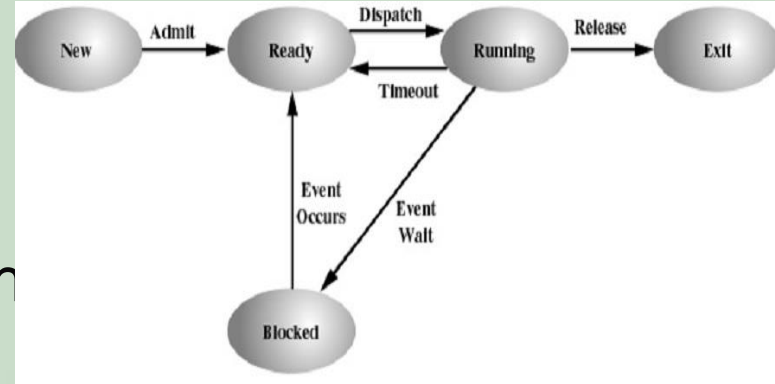


A Five-State Model

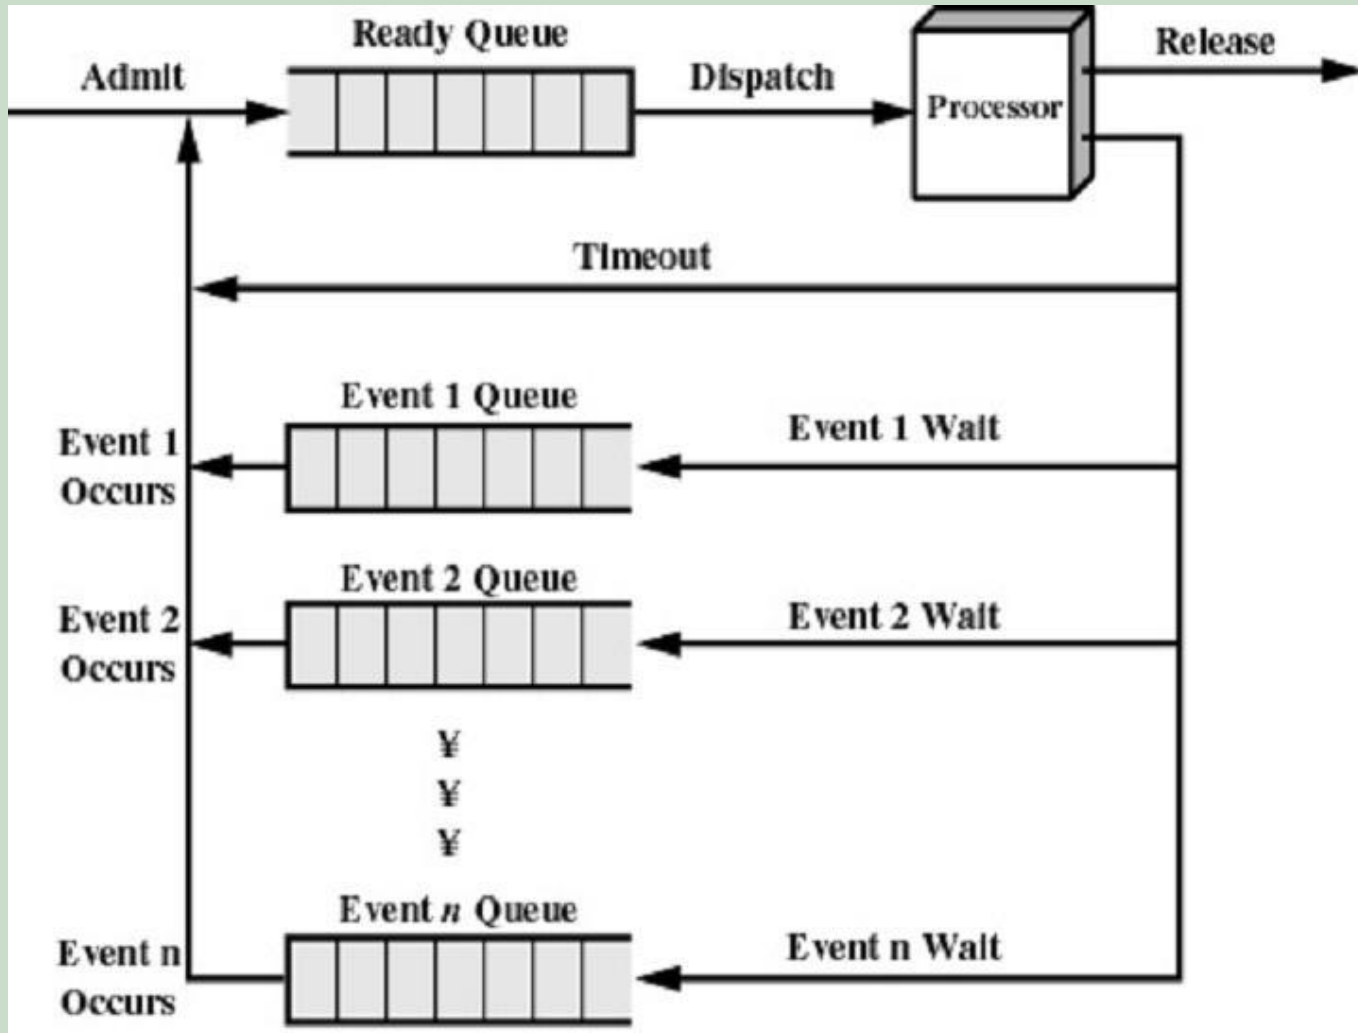


Five-state Model Transitions

- ↻ **Null → New:** Process is created
- ↻ **New → Ready:** O.S. is ready to handle another process
- ↻ **Ready → Running:** Select another process to run
- ↻ **Running → Exit:** Process has terminated
- ↻ **Running → Ready:** End of time slice or higher-priority process is ready
- ↻ **Running → Blocked:** Process is waiting for an event
- ↻ **Blocked → Ready:** The event a process is waiting for has occurred, can continue
- ↻ **Ready → Exit:** Process terminated by O.S. or parent
- ↻ **Blocked → Exit:** Same reasons



Multiple Blocked Queues




Suspended Processes

- Processor is faster than I/O so all processes could be waiting for I/O
- Swap these processes to disk to free up more memory
- Blocked state becomes suspended state when swapped to disk
- Two new states
 - ∞ Blocked, suspend
 - ∞ Ready, suspend

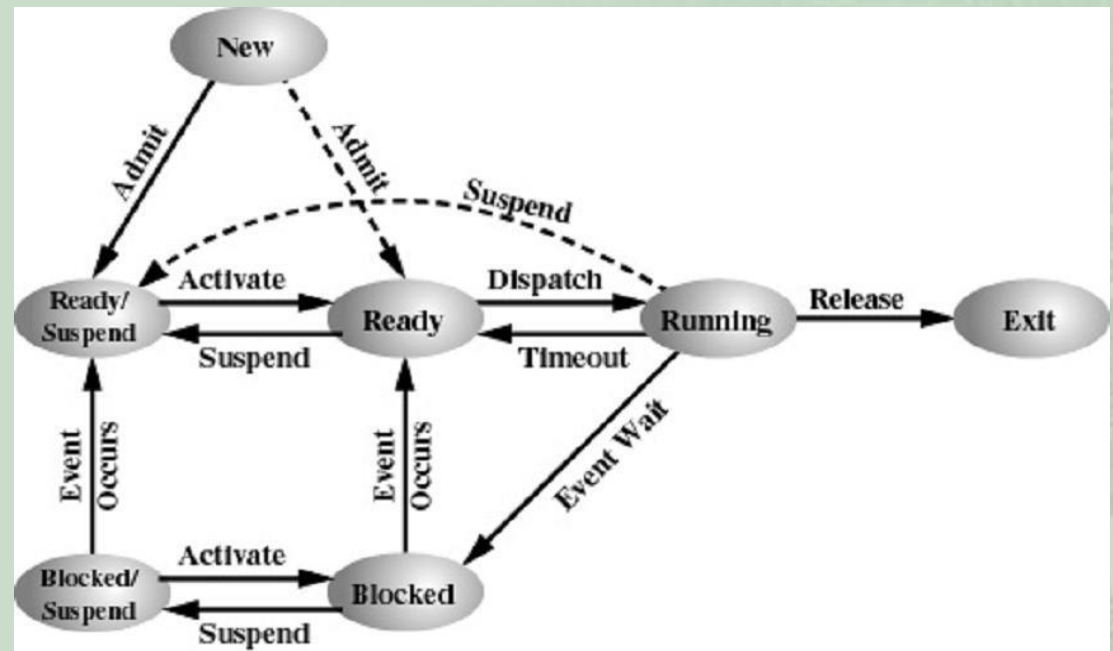


Suspended State Scheduling



One Suspend State

Two Suspend State



Reasons for Process Suspension

- Swapping
 - ☞ The OS needs to release sufficient main memory to bring in a process that is ready to execute
- Other OS reason
 - ☞ The OS may suspend a background or utility process or a process that is suspected of causing a problem
- Interactive user request
 - ☞ A user may wish to suspend execution of a program for purposes of debugging or in connection with the use of a resource
- Timing
 - ☞ A process may be executed periodically and may be suspended while waiting for the next time interval
- Parent process request
 - ☞ A parent process may wish to suspend execution of a descendent to examine or modify the suspended process

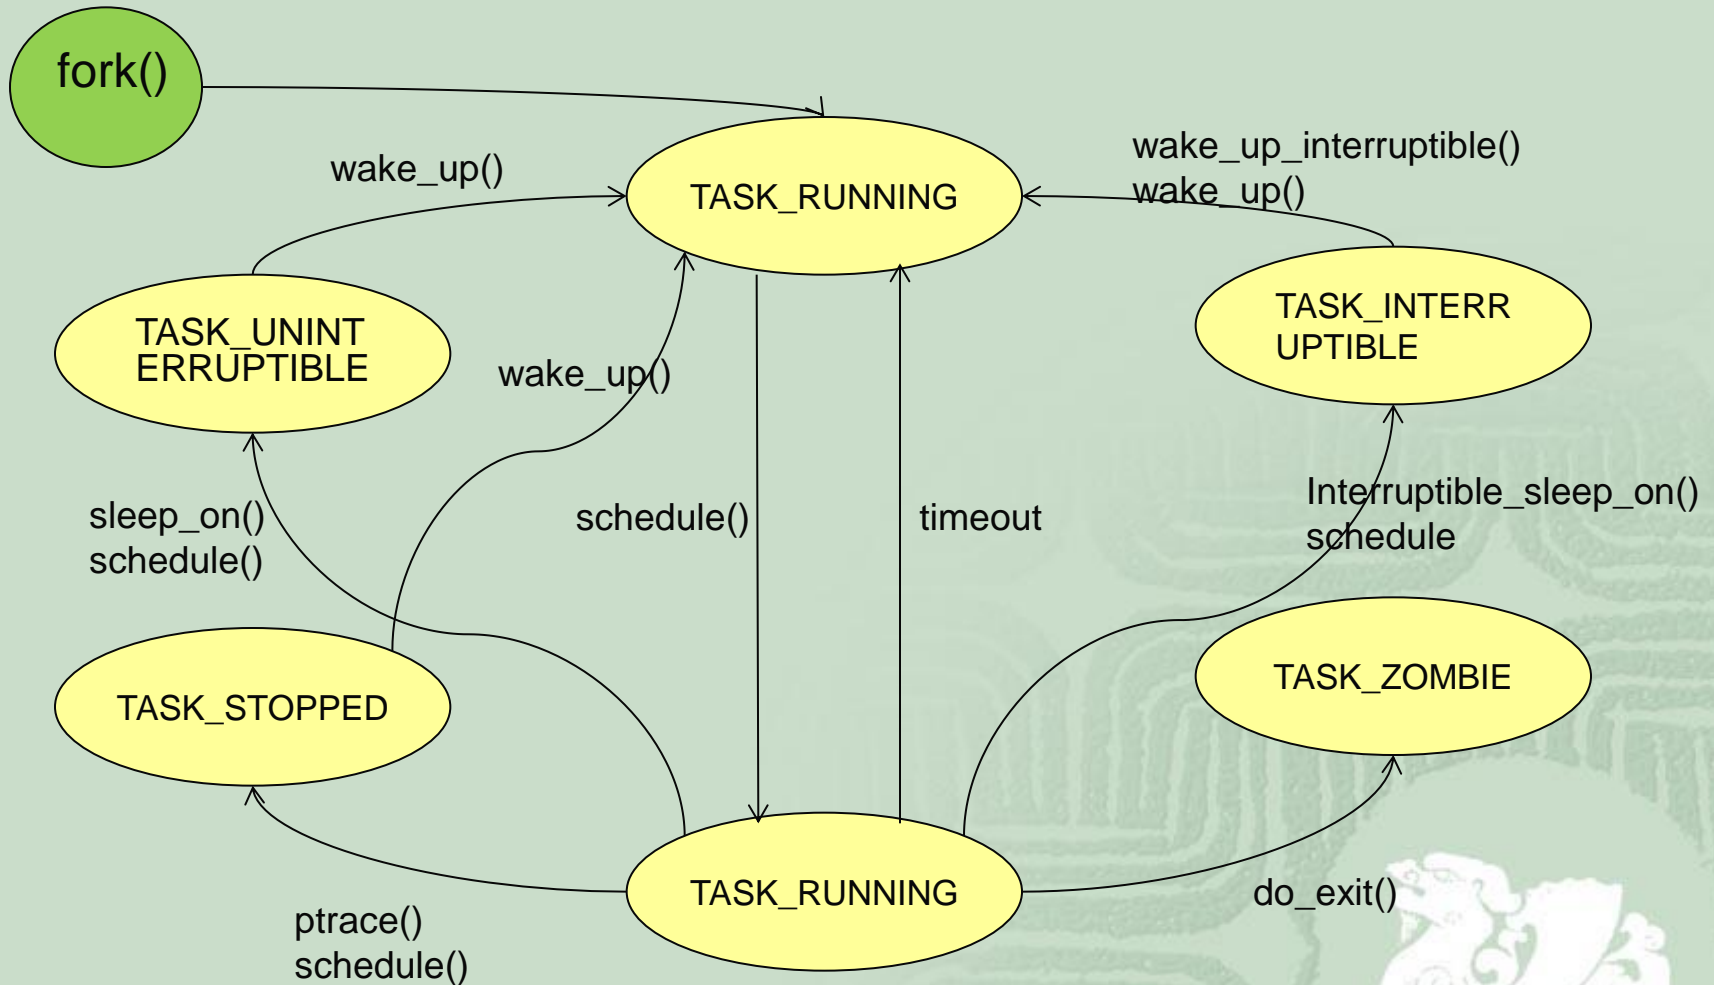


The 'states' of a Linux process

- A process can be in one of several states:
 - ❧ 0: TASK_RUNNING (R)
 - ❧ 1: TASK_INTERRUPTIBLE (S)
 - ❧ 2: TASK_UNINTERRUPTIBLE (D)
 - ❧ 4: TASK_ZOMBIE (Z)
 - ❧ 8: TASK_STOPPED (T)



State Transitions of a Linux Process



Process Concepts

- Process Definition
- Process Creation
- Process States and State Transitions
- Important Information Associated with Processes

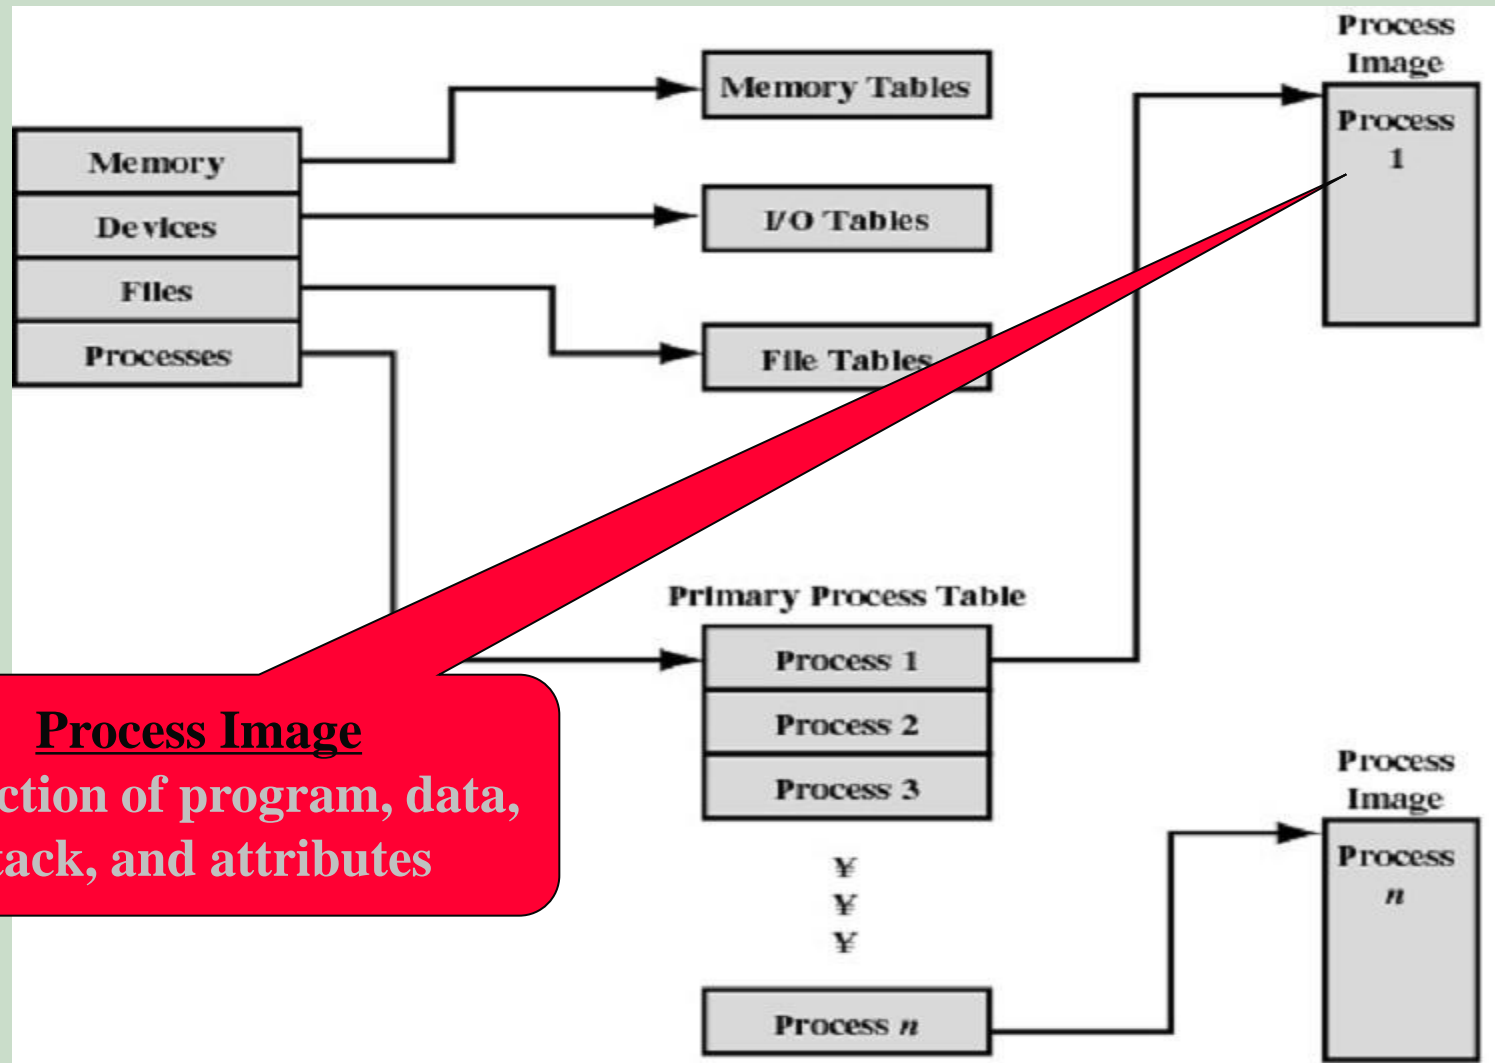


Operating System Control Structures

- Information the OS needs to keep around:
 - ☞ Memory tables
 - ☞ I/O tables
 - ☞ File tables
 - ☞ Process tables
- Tables are constructed for each entity the operating system manages
- The OS has to manage and protect these tables



Process Control Structures



Control Tables

- Memory Tables
 - ❧ Allocation of main memory to processes
 - ❧ Allocation of secondary memory to processes
 - ❧ Protection attributes for access to shared memory regions
 - ❧ Information needed to manage virtual memory
- I/O device is available or assigned
 - ❧ Status of I/O operation
 - ❧ Location in main memory being used as the source or destination of the I/O transfer

Control Tables (continued...)

■ File Tables

- ∞ Existence of files
- ∞ Location on secondary memory
- ∞ Current Status
- ∞ Attributes
- ∞ Usually maintained by a file management system

■ Process Table

- ∞ Process ID
- ∞ Process state
- ∞ Location in memory



Process Location

- A Process includes:
 - ❧ Data locations for local and global variables
 - ❧ Any defined constants
 - ❧ Stack
 - ❧ Heap
 - ❧ Set of programs to be executed
- Process Control Block (PCB)
 - ❧ Collection of attributes



Process Control Block (PCB)

pointer	process state
process number	
program counter	
registers	
memory limits	
list of open files	
⋮	

Process Control Block

- Process identification
- Processor State Information
 - ❧ User-Visible Registers
 - ❧ Control and Status Registers
 - ❧ Stack Pointers
- Process Control Information
 - ❧ Scheduling and State Information
 - *Process state*
 - *Priority*
 - *Scheduling-related information*
 - *Event*
 - ❧ Data Structuring
 - ❧ Interprocess Communication
 - ❧ Process Privileges
 - ❧ Memory Management
 - ❧ Resource Ownership and Utilization



Process Control Block (continued...)

- Process identification

- ⌘ Numeric identifiers that may be stored with the process control block and include

- Identifier of this process
 - Identifier of creator process (parent process)
 - User identifier

- Processor State Information

- ⌘ User-Visible Registers

- ⌘ Control and Status Registers

- ⌘ Stack Pointers



Process Control Block (continued...)

■ Processor State Information

⌘ User-Visible Registers

- A register is one that may be referenced by means of the machine language that the processor executes. (Typically, 8 to 32)

⌘ Control and Status Registers

- *Program counter*: Address of the next instruction to be fetched
- • *Condition codes*: Results of the most recent arithmetic or logical operation (e.g., sign, zero, carry, equal, overflow)
- *Status information*: Interrupt enabled/disabled flags, execution mode

⌘ Stack Pointers

- Each process has one or more last-in-first-out (LIFO) system stacks associated with it. A stack is used to store parameters and calling addresses for procedure and system calls. The stack pointer points to the top of the stack.

Several dozen fields

- Dozens of separate items of information are kept in a Linux 'task_struct'; e.g.:

- ❧ pid; // process ID-number
- ❧ state; // current task-state
- ❧ priority; // current task-priority
- ❧ start_time; // time when task was created
- ❧ sleep_time; // time when task began sleeping
- ❧ . . . // ... many more items ...

- This info is used by the Linux 'scheduler'



Execution Stack in C

- The C runtime system keeps track of the chain of active functions with a stack
- When a function is called, the runtime system pushes on the stack a frame containing
 - ∞ Local variables and return value
 - ∞ Program counter, keeping track of the statement being executed
- When a function ends, its frame is popped from the stack and control is passed to the function on top of the stack

```
main() {  
    int i = 5;  
    foo(i);  
}
```

```
foo(int j) {  
    int k;  
    k = j+1;  
    bar(k);  
}
```

```
bar(int m) {  
    ...  
}
```

