



# Threads

# Consider

- A Process is a unit of:
  - ✧ Resource ownership
    - Address space
    - I/O channels, devices, files
  - ✧ Execution path
    - Interleaved with other processes
    - State
- What if we treat each independently?
  - ✧ Unit of resource ownership → process
  - ✧ Unit of execution → thread



# Process Context Switch

## Process Context switch:

allocate CPU from one process to another.

## A Process context includes two portions:

CPU context and Storage context.


- CPU context: program counter, registers, stack/heap pointers and other control registers. **Easy to switch.**
- Storage context: program code, data, address space, memory mapping, (disk) swapping, resources, etc. **Hard and time consuming to switch.**




# Multi-threading

- Operating system supports multiple threads of execution within a single process.
  - ✧ Traditional approach is single-threaded.
- Examples:
  - ✧ MS-DOS supports a single user process and a single thread.
  - ✧ UNIX supports multiple user processes but only supports one thread per process.
  - ✧ Java run-time environment is a system of one process with multiple threads.
  - ✧ Windows 2000 (W2K), Solaris, Linux, Mach, and OS/2 support multiple processes, each of which supports multiple threads.



# Threads and Processes





**one process  
one thread**



**one process  
multiple threads**



**multiple processes  
one thread per process**



**multiple processes  
multiple threads per process**



# Threads

- One view of a thread is as an independent program counter operating within a process.
- A thread consists of:
  - ∞ a thread execution state (Running, Ready, etc.)
  - ∞ a CPU context (program counter, register set.)
  - ∞ an execution stack.
  - ∞ access to the memory and resources of its process (shared with all other threads in that process.)






# Threads (continued...)


- Thus, all of the threads of a process share the state and resources of the parent process (memory space and code section.)
- A process is defined to have at least one thread of execution (the process itself) and may launch other threads which execute concurrently with the process.
- Key benefits:
  - ✧ Far less time to create/terminate.
  - ✧ Switching between threads is faster.
  - ✧ No memory management issues, etc.
  - ✧ Can enhance communication efficiency.



# Threads (continued...)



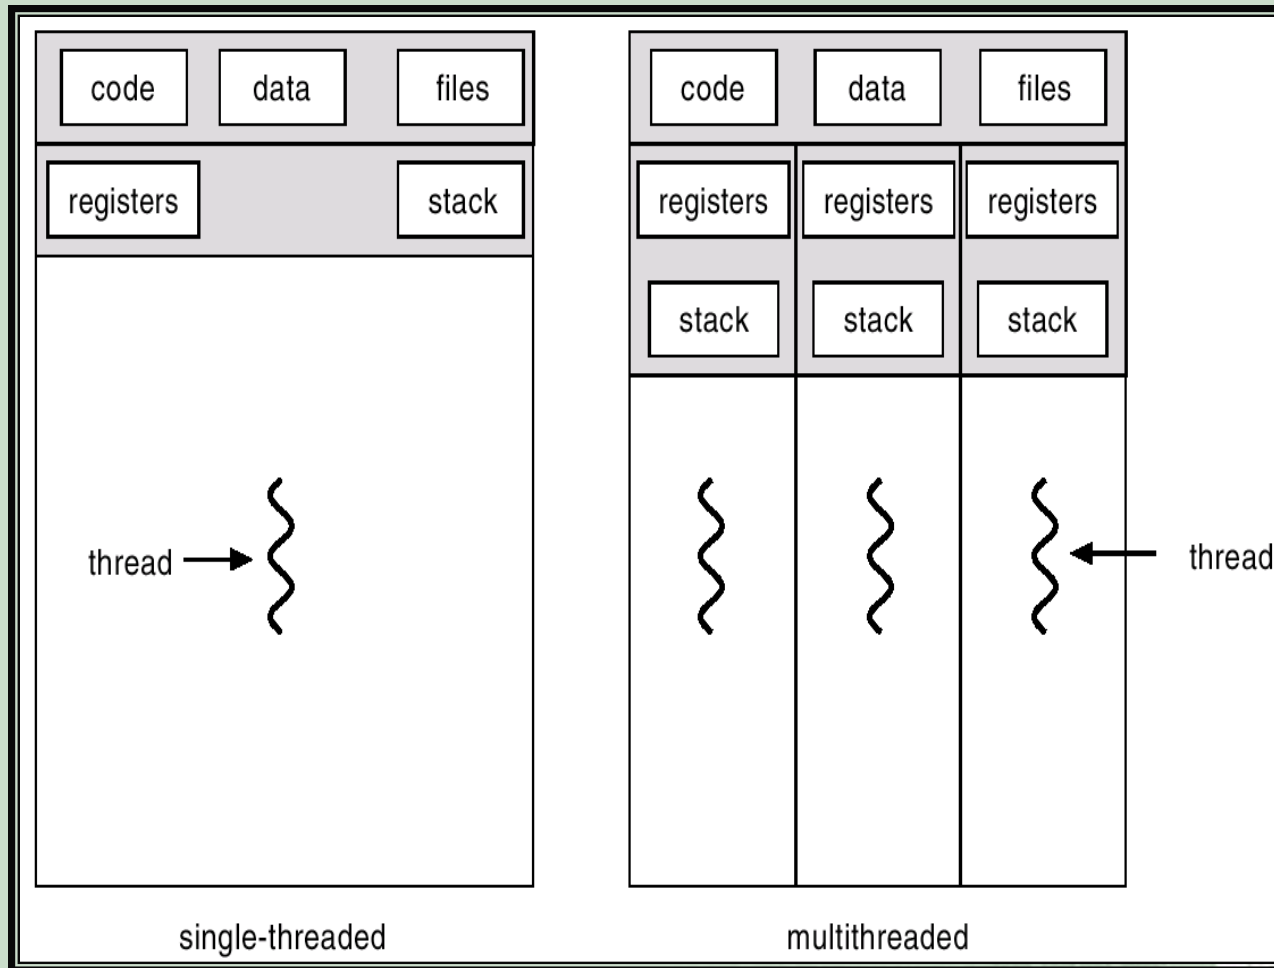
(a) Task graph of a program



(b) Thread structure of a task



# Single and Multithreaded Processes




# Using Threads




- Multiple threads in a single process
  - ✧ Separate control blocks for the process and each thread
  - ✧ Can quickly switch between threads
  - ✧ Can communicate without invoking the kernel
- Examples
  - ✧ Producer/Consumer
  - ✧ Concurrent services
  - ✧ Faster Execution – Read one set of data while processing another set
  - ✧ GUI and worker



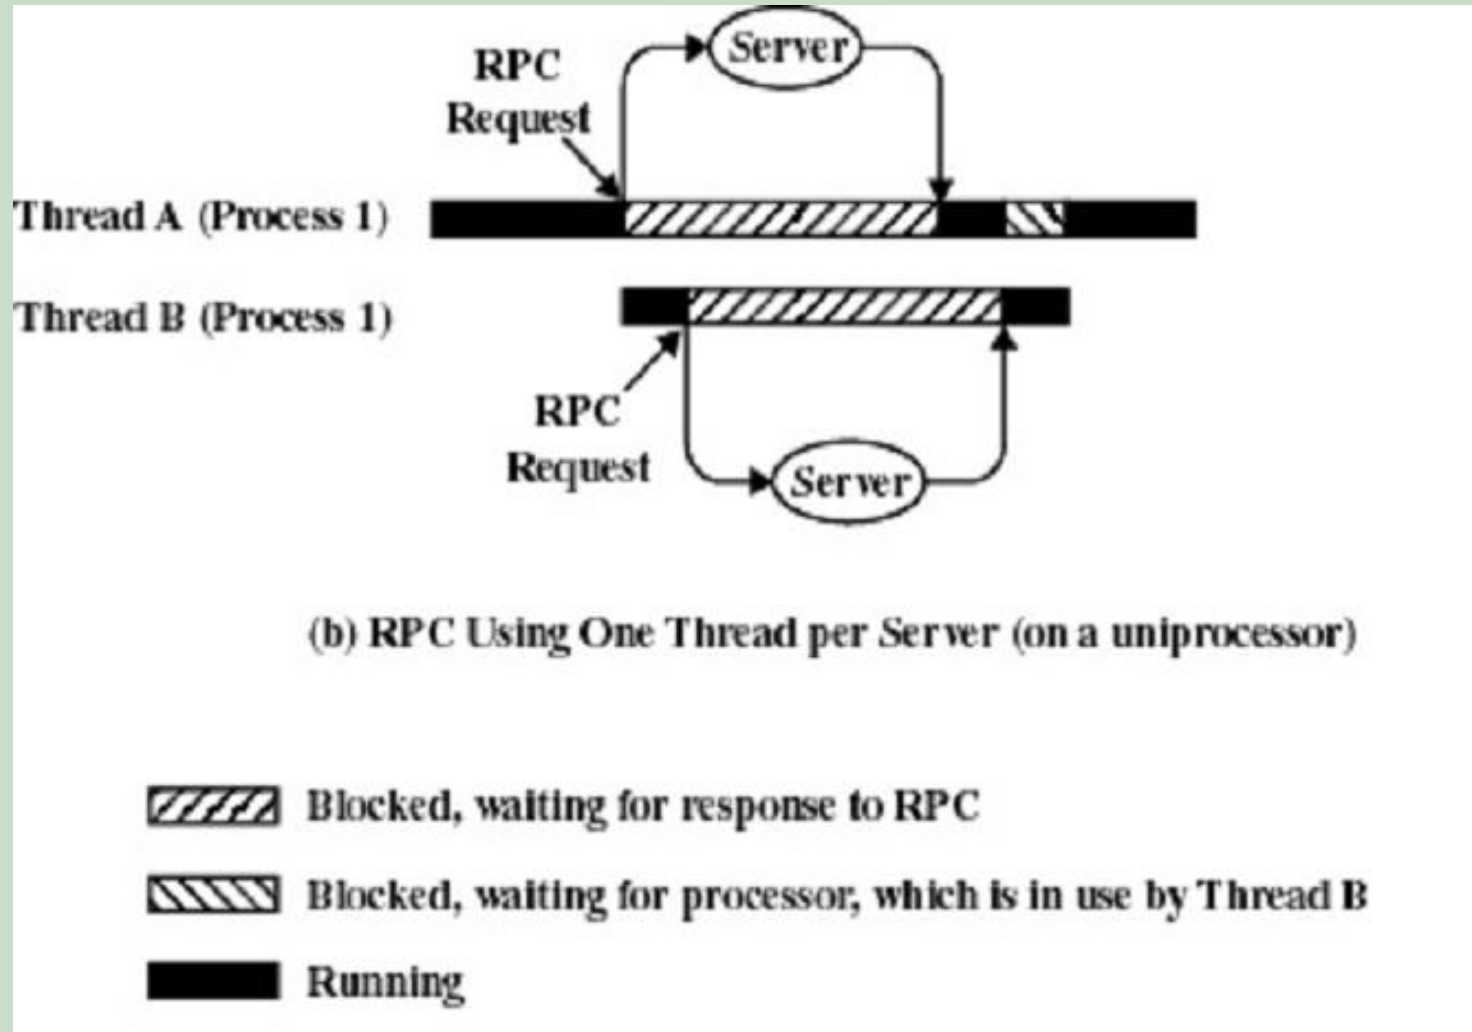
# Remote Procedure Calls




(a) RPC Using Single Thread

-  Blocked, waiting for response to RPC
-  Blocked, waiting for processor, which is in use by Thread B
-  Running

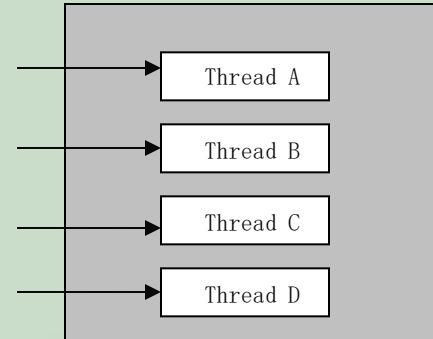
# Remote Procedure Calls Using Threads




## Design of Concurrent Server




(a) Center distributor



(b) Concurrent threads



(c) Center scheduler



(d) Round-robin schedule

# Thread States

- **Thread operations**

- ☞ *Spawn* – Creating a new thread
- ☞ *Block* – Waiting for an event
- ☞ *Unblock* – Event happened, start new
- ☞ *Finish* – This thread is completed

- **Generally a thread can block without blocking the remaining threads in the process**





# Thread issues

- How should threads be scheduled compared to processes?
  - ✧ Equal to processes
  - ✧ Within the parent processes quantum
- How are threads implemented?
  - ✧ kernel support (system calls)
  - ✧ user level threads



# User-Level Threads

- All thread management is done by the application
- The kernel is not aware of the existence of threads
- Thread switching does not require kernel mode privileges
- Scheduling is application specific

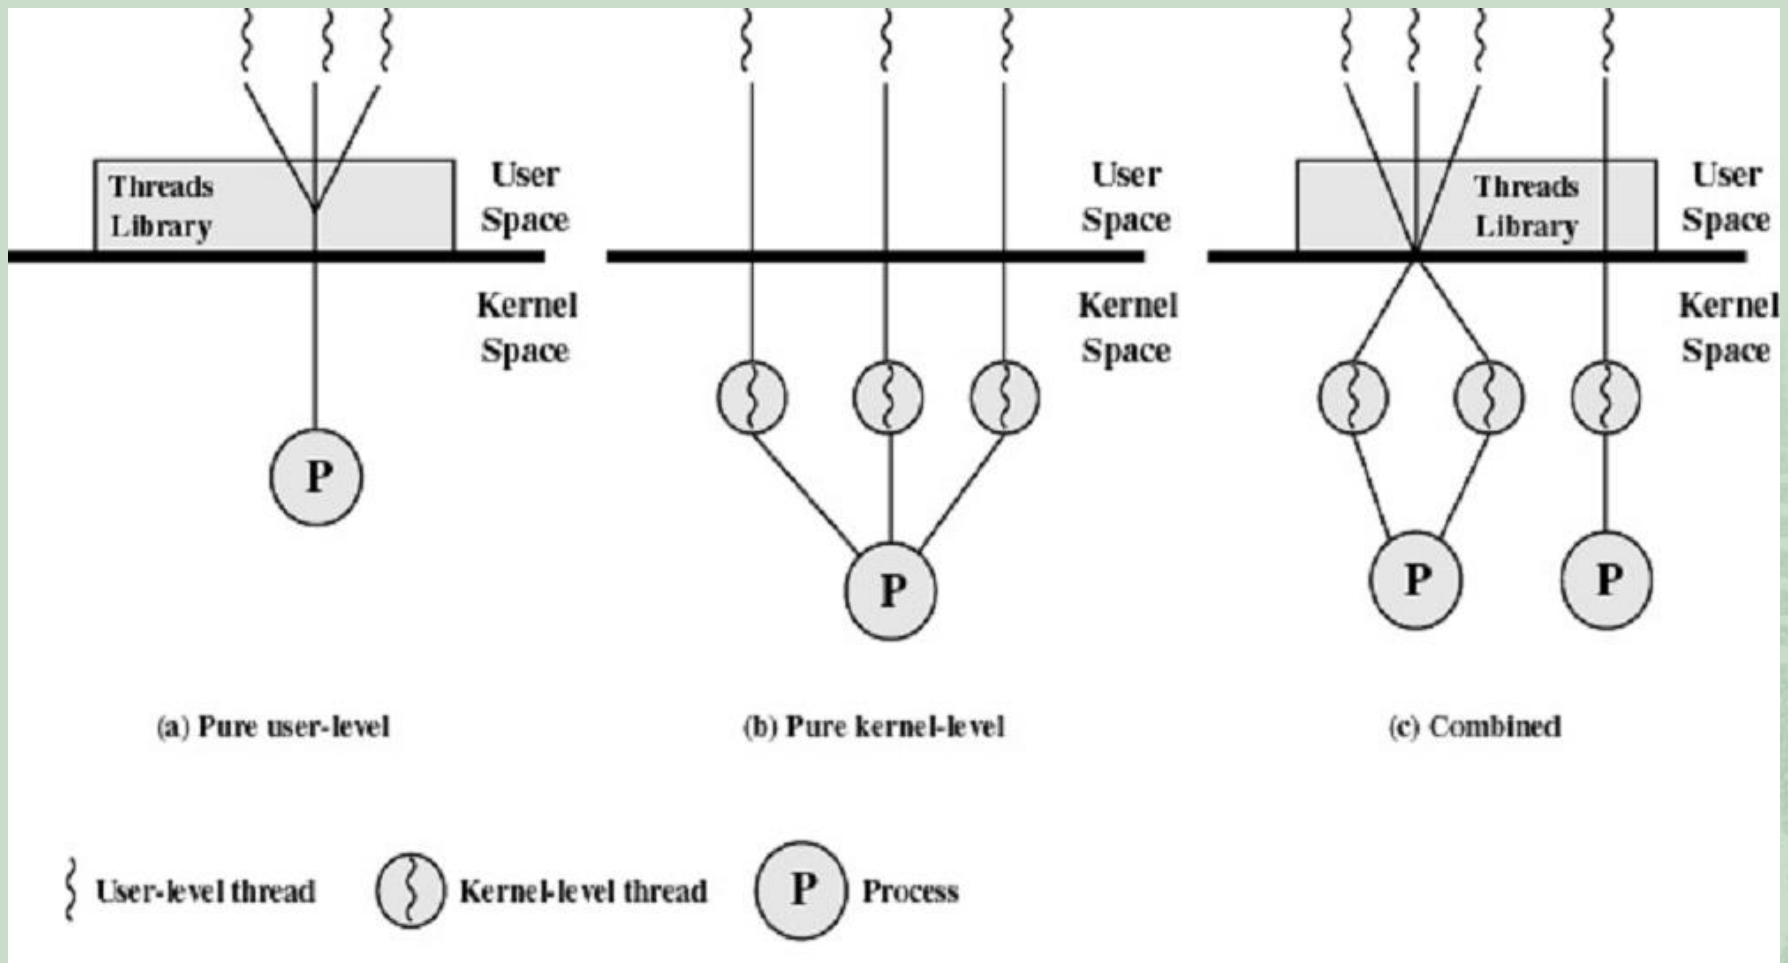


# Kernel-Level Threads


- Kernel maintains context information for the process and the threads
- Switching between threads requires the kernel
- Windows NT and OS/2 are examples of this approach



# User-Level and Kernel-Level Threads



# Many-to-One Model




# Many-to-One

- Many user-level threads mapped to single kernel thread.
- Used on systems that do not support kernel threads.
- Example: GNU Portable Threads





# One-to-one Model




# One-to-One

- Each user-level thread maps to kernel thread.
- Examples
  - Windows 32
  - old Linux, Solaris



# Many-to-Many Model



# Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads.
- Allows the operating system to create a sufficient number of kernel threads.
- Solaris 2
- Windows 7

