


CPU Scheduling

Round Robin (RR)

- Each process gets a small amount of CPU time (*time quantum*), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n-1)q$ time units.
- Performance
 - ⌘ q large \Rightarrow FCFS
 - ⌘ q small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high.


Round Robin (continued...)



Process	Arrival	Service
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



Round Robin (continued...)




Process	Arrival	Service
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2

Suppose $q = 1$;

Turnaround time: p1: $4 - 0 = 4$; p2: $18 - 2 = 16$; p3: $17 - 4 = 13$;

P4: $20 - 6 = 14$; p5: $15 - 8 = 7$;


Ave TAT: $(4 + 16 + 13 + 14 + 7) / 5 = 8.8$



Example of RR with Time Quantum = 20

<u>Process</u>	<u>Burst Time</u>
P_1	53
P_2	17
P_3	68
P_4	24

- The Gantt chart is:




Round Robin (continued...)

- Particularly effective in time-sharing or transaction processing system.
- Generally favors CPU bound processes as I/O bound give up their time slice while CPU bound use their complete time quantum.
- Could be improved using a virtual round robin (VRR) scheduling scheme – implement an auxiliary I/O queue which gets preference over the main queue.




Time Quantum and Context Switch Time



Shortest Process Next (SPN)

Process	Arrival	Service
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2




- Shortest Process Next - Non-preemptive policy
- Process with shortest expected processing time is selected next
- Short process jumps ahead of longer processes

Example of Non-Preemptive SPN

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

- SPN (non-preemptive)



- Average waiting time = $(0 + 6 + 3 + 7)/4 = 4$


Shortest Process Next (continued...)

- May be impossible to know or at least estimate the required processing time of a process.
 - ❧ For batch jobs, require a programmer's estimate.
 - If estimate is substantially off, system may abort job.
 - ❧ In a production environment, the same jobs run frequently and statistics may be gathered.
 - ❧ In an interactive environment, the operating system may keep a running average of each “burst” for each process.
- SPN could result in starvation of longer processes if there is a steady supply of short processes.
- Not suitable for time-sharing or transaction processing environment because of lack of preemption.



Process	Arrival	Service
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2

Shortest Remaining Time (SRT)




- The shortest remaining time (SRT) policy is a preemptive version of SPN.
- Must estimate expected remaining processing time

Example of Preemptive SRT

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

- SRT (preemptive)



- Average waiting time = $(9 + 1 + 0 + 2)/4 = 3$



Shortest Remaining Time

(continued...)

- When a new process joins the ready queue, it may have a shorter remaining time and preempts the current process.
- SRT does not bias in favor of long processes (as does FCFS)
- Unlike RR, no additional interrupts are generated reducing overhead.
- Superior turnaround performance to SPN, because a short job is given immediate preference to a running longer job.