



Mutual Exclusion

Hardware Solutions

- Interrupt disabling
 - ∞ not all interrupts need to be disabled
 - ∞ might not work in multiprocessor environment
- Special machine instructions
 - ∞ test and set instruction (**TSET**)
 - ∞ exchange instruction (**XCHG**)
 - ∞ advantages
 - works on any number of processors sharing memory
 - simple and supports multiple critical sections
 - ∞ disadvantages
 - busy waiting is employed
 - starvation and deadlock are possible



Test-and-Set

```
bool testset(int &i)
{  if (*i == 0)
    {  *i = 1;
        return false;
    }
    else return true;
}
```

An atomic action
An Instruction

- Shared variable b (initialized to 0)
- Only the first P_i who sets b enters CS

```
int b;
```

```
testset(&b)
```



Mutual Exclusion with Test-and-Set

- Shared data:

boolean lock = false;

- Process P_i

do {

 while (testset(&lock)) ; // busy waiting

 critical section

 *lock = false;

 remainder section

}



Synchronization Hardware

- Atomically swap two variables.

```
void swap(boolean &a, boolean &b)
{
    boolean temp = *a;
    *a = *b;
    *b = temp;
}
```



Mutual Exclusion with Swap

- Shared data (initialized to **false**):
boolean lock;
- Process P_i
do {
 key = true;
 while (key == true) swap(&lock,&key);
 critical section
 lock = false;
 remainder section
}



Concurrency Control

Mutual Exclusion and Synchronization

Principles of Concurrency

Mutual Exclusion

Software Solutions

Hardware Solutions

☞ Semaphores

Producer/Consumer

Bounded Buffer

Message Passing

Monitors



Mutual Exclusion Issues

- Software solutions are complex
 - ☞ must be used when no Hardware support
- Hardware instructions work
 - ☞ Still kind of complex
- Both make the processes **busy-wait**
 - ☞ Wastes CPU cycles




Semaphores

- A synchronization tool provided by the OS.
- Two or more processes can cooperate by using simple signals – a process can be forced to stop until it has received a specific signal.
- Semaphore operations:
 - ⌘ initialize to a non-negative value
 - ⌘ *wait* decrements the semaphore value – if negative, the process is blocked
 - ⌘ *signal* increments the semaphore value – if not positive, then a process blocked by a *wait* operation is unblocked.

Semaphores

- Modifications to semaphore variable are atomic.
- Two forms of semaphores:
 - ☞ binary semaphore
 - easier to implement
 - also called a mutex
 - ☞ counting semaphore
 - *wait* - # of processes waiting on the resource
 - *signal* - # of resources available
- Blocked processes are held in waiting queues
 - ☞ strong semaphores unblock the longest waiting process
 - guarantee freedom from starvation
 - ☞ weak semaphores unblock without regard to order



The Critical Section Problem...

Shared semaphore: mutex = 1;

repeat

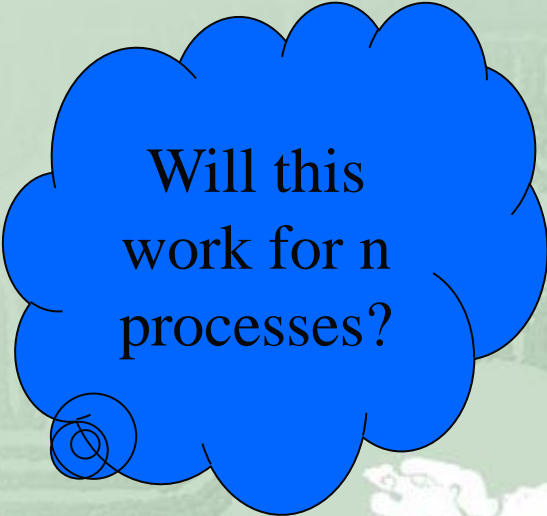
wait(mutex)

critical section

signal(mutex)

remainder section

until false



Will this
work for n
processes?



wait/signal

```
struct semaphore
{  int value;
   Queue processes;
}

wait(Semaphore s)
{  s.value = s.value - 1;
   if (s.value < 0)
   {  add this process to s.queue
      block; // the caller is BLOCKED
   }
}

signal(Semaphore s)
{  s.value = s.value + 1;
   if (s.value <= 0)
   {  remove a process P from s.queue
      wakeup(P); // make P READY
   }
}
```



Autonomy

- Critical
 - ☞ Semaphore operations must be atomic
- Uniprocessor
 - ☞ simply inhibit interrupts (normal user can't)
 - ☞ Use **testset** to create a mutex in the calls
- Multiprocessor
 - ☞ hardware must provide special support, or
 - ☞ use software solutions

```
wait(Semaphore s)
{  while (testset(&lock)) ;
    s.value--;
    if (s.value < 0)
    {  add process to s.queue
       lock = FALSE;
       block;
    }
    else lock = FALSE;
}
```



Autonomy

- Commonly used ME mechanisms are:
- Semaphore, lock, shared variables, or conditional variables
- P/V:
 $P(s)$ [while $(s==0)$ wait; $s = s-1$]
 $V(s)$ [$s = s+1$]
- lock/unlock
- wait/signal
- monitors, etc.



Consider...

semaphore S, Q; // initial value 1

P0:

P(S);

P(Q);

.

.

.

V(S);

V(Q);

P1:

P(Q);

P(S);

.

.

.

V(Q);

V(S);

Is there anything wrong with this?



A simple example

Expression: $(x^2) / (y-z) + \sin(x)$

Shared variables: t_1, t_2, t_3, t_4, t_5 ;


P1: $t_1 = x^2$;

P2: $t_2 = y - z$;

P3: $t_3 = \sin(X)$;

P4: $t_4 = t_1/t_2$;

P5: $t_5 = t_4 + t_3$



Concurrency control

semaphore s1, s2, s3, s4; // initial 0s


P1: $t1 = x^2$; $V(s1)$;

P2: $t2 = y - z$; $V(s2)$;

P3: $t3 = \sin(X)$; $V(s3)$;

P4: $P(s1)$; $P(s2)$; $t4 = t1/t2$; $V(s4)$;

P5: $P(s3)$; $P(s4)$; $t5 = t4 + t3$;



The Producer-Consumer Problem

Producer/Consumer

```
int counter = 0;  
Item buffer[n];
```

Producer

```
while(1) {  
    ...  
    produce an item in nextp  
    ...  
    while(counter == n)  
        do no-op  
    buffer[in] = nextp  
    in = (in + 1) mod n  
    counter = counter + 1  
}
```

Consumer

```
while(1) {  
    ...  
    while(counter == 0)  
        do no-op  
    nextc = buffer[out]  
    out = (out + 1) mod n  
    counter = counter - 1  
    ...  
    consume the item in nextc  
    ...  
}
```



Bounded Buffer Solution

Shared semaphore: `empty = n, full = 0;`
Item buffer[n];

```
while(1){  
    produce an item in nextp  
  
    wait(empty);  
  
    buffer[in] = nextp  
    in = (in + 1) mod n  
  
    signal(full);  
  
}
```

Producer

```
while(1){  
    wait(full);  
  
    nextc = buffer[out]  
    out = (out + 1) mod n  
  
    signal(empty);  
  
    consume the item in nextc  
}
```

Consumer



Any problem?

Bounded Buffer Solution

Shared semaphore: `empty = n, full = 0, mutex = 1;`
Item `buffer[n];`

```
while(1){  
    produce an item in nextp  
  
    wait(empty);  
    wait(mutex);  
  
    buffer[in] = nextp  
    in = (in + 1) mod n  
  
    signal(mutex);  
    signal(full);  
}
```

Producer

```
while(1){  
    wait(full);  
    wait(mutex);  
  
    nextc = buffer[out]  
    out = (out + 1) mod n  
  
    signal(mutex);  
    signal(empty);  
  
    consume the item in nextc  
}
```

Consumer



Mutex + Synchronization: multiple producers and comusers

Message Passing

- A general method used for interprocess communication (IPC)
 - ✧ for processes inside the same computer
 - ✧ for processes in a distributed system
- Another means to provide process synchronization and mutual exclusion
- We have at least two primitives:
 - ✧ `send(destination, message)`
 - ✧ `receive(source, message)`
- May or may not be blocking



Synchronization

- For the sender: it is more natural not to be blocked
 - ⌘ can send several messages to multiple destinations
 - ⌘ sender usually expects acknowledgment of message receipt (in case receiver fails)
 - ⌘ *PostMessage()* is asynchronous – returns immediately
 - ⌘ *SendMessage()* is synchronous – block until message delivered and processed
- For the receiver: it is more natural to be blocked after issuing *ReceiveMessage()*
 - ⌘ the receiver usually needs the info before proceeding
 - ⌘ but could be blocked indefinitely if sender process fails before sending reply

Addressing in message passing

- Direct addressing:

- ✧ when a specific process identifier is used for source/destination
- ✧ but it might be impossible to specify the source ahead of time (ex: a print server)


- Indirect addressing (more convenient):

- ✧ messages are sent to a shared **mailbox** which consists of a queue of messages
- ✧ senders place messages in the mailbox, receivers pick them up



Mailboxes and Ports

- A mailbox can be private
 - ⌘ one sender/receiver pair
- A mailbox can be shared among several senders and receivers
 - ⌘ OS may then allow the use of message types (for selection)
- **Port:** a mailbox associated with one receiver and multiple senders
 - ⌘ used for client/server application: the receiver is the server



Ownership of ports and mailboxes

- A port is usually owned and created by the receiving process
- The port is destroyed when the receiver terminates
- The OS creates a mailbox on behalf of a process (which becomes the owner)
- The mailbox is destroyed at the owner's request or when the owner terminates



Message format

- Consists of header and body of message
- Control information:
 - ⌘ what to do if run out of buffer space
 - ⌘ sequence numbers
 - ⌘ priority...
- **Queuing discipline: usually FIFO but can also include priorities**

