



Mutual Exclusion

The Producer-Consumer Problem

Producer/Consumer

```
int counter = 0;  
Item buffer[n];
```

Producer

```
while(1) {  
    ...  
    produce an item in nextp  
    ...  
    while(counter == n)  
        do no-op  
    buffer[in] = nextp  
    in = (in + 1) mod n  
    counter = counter + 1  
}
```

Consumer

```
while(1) {  
    ...  
    while(counter == 0)  
        do no-op  
    nextc = buffer[out]  
    out = (out + 1) mod n  
    counter = counter - 1  
    ...  
    consume the item in nextc  
    ...  
}
```



Bounded Buffer Solution

Shared semaphore: `empty = n, full = 0;`
Item buffer[n];

```
while(1){  
  
    produce an item in nextp  
  
    wait(empty);  
  
    buffer[in] = nextp  
    in = (in + 1) mod n  
  
    signal(full);  
  
}
```

Producer

```
while(1){  
  
    wait(full);  
  
    nextc = buffer[out]  
    out = (out + 1) mod n  
  
    signal(empty);  
  
    consume the item in nextc  
  
}
```

Consumer



Any problem?

Bounded Buffer Solution

Shared semaphore: `empty = n, full = 0, mutex = 1;`
Item buffer[n]; `int in = out = 0;`

```
while(1){  
    produce an item in nextp  
  
    wait(empty);  
    wait(mutex);  
  
    buffer[in] = nextp  
    in = (in + 1) mod n  
  
    signal(mutex);  
    signal(full);  
}
```

Producer

```
while(1){  
    wait(full);  
    wait(mutex);  
  
    nextc = buffer[out]  
    out = (out + 1) mod n  
  
    signal(mutex);  
    signal(empty);  
  
    consume the item in nextc  
}
```

Consumer



Mutex + Synchronization: multiple producers and consumers

Message Passing

- A general method used for interprocess communication (IPC)
 - ✧ for processes inside the same computer
 - ✧ for processes in a distributed system
- Another means to provide process synchronization and mutual exclusion
- We have at least two primitives:
 - ✧ `send(destination, message)`
 - ✧ `receive(source, message)`
- May or may not be blocking



Synchronization

- For the sender: it is more natural not to be blocked
 - ⌘ can send several messages to multiple destinations
 - ⌘ sender usually expects acknowledgment of message receipt (in case receiver fails)
- For the receiver: it is more natural to be blocked after issuing *ReceiveMessage()*
 - ⌘ the receiver usually needs the info before proceeding
 - ⌘ but could be blocked indefinitely if there is no sender



Addressing in message passing

- Direct addressing:

- ✧ when a specific process identifier is used for source/destination
- ✧ but it might be impossible to specify the source ahead of time (ex: a print server)


- Indirect addressing (more convenient):

- ✧ messages are sent to a shared **mailbox** which consists of a queue of messages
- ✧ senders place messages in the mailbox, receivers pick them up



Mailboxes and Ports

- A mailbox can be private
 - ⌘ one sender/receiver pair
- A mailbox can be shared among several senders and receivers
 - ⌘ OS may then allow the use of message types (for selection)
- **Port:** a mailbox associated with one receiver and multiple senders
 - ⌘ used for client/server application: the receiver is the server



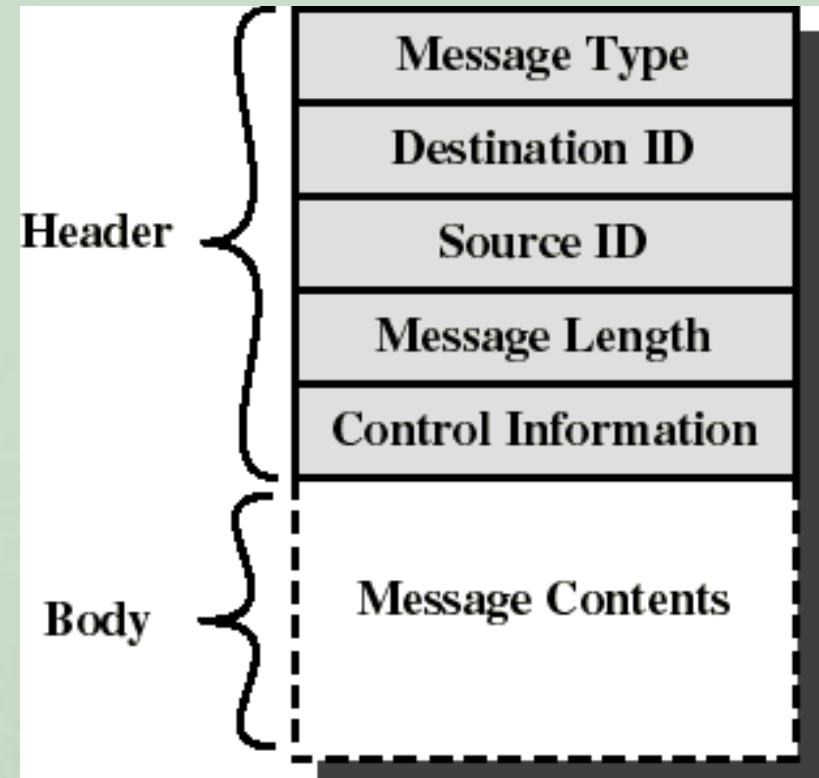
Ownership of ports and mailboxes

- A port is usually owned and created by the receiving process
- The port is destroyed when the receiver terminates
- The OS creates a mailbox on behalf of a process (which becomes the owner)
- The mailbox is destroyed at the owner's request or when the owner terminates




Message format

- Consists of header and body of message
- Control information:
 - ⌘ what to do if run out of buffer space
 - ⌘ sequence numbers
 - ⌘ priority...
- **Queuing discipline: usually FIFO but can also include priorities**



Monitor

- A software module containing:
 - ⌘ one or more procedures
 - ⌘ an initialization sequence
 - ⌘ local data variables
- Characteristics:
 - ⌘ local variables accessible only by monitor's procedures
 - ⌘ a process enters the monitor by invoking one of its procedures
 - ⌘ only one process can be in the monitor at any one time



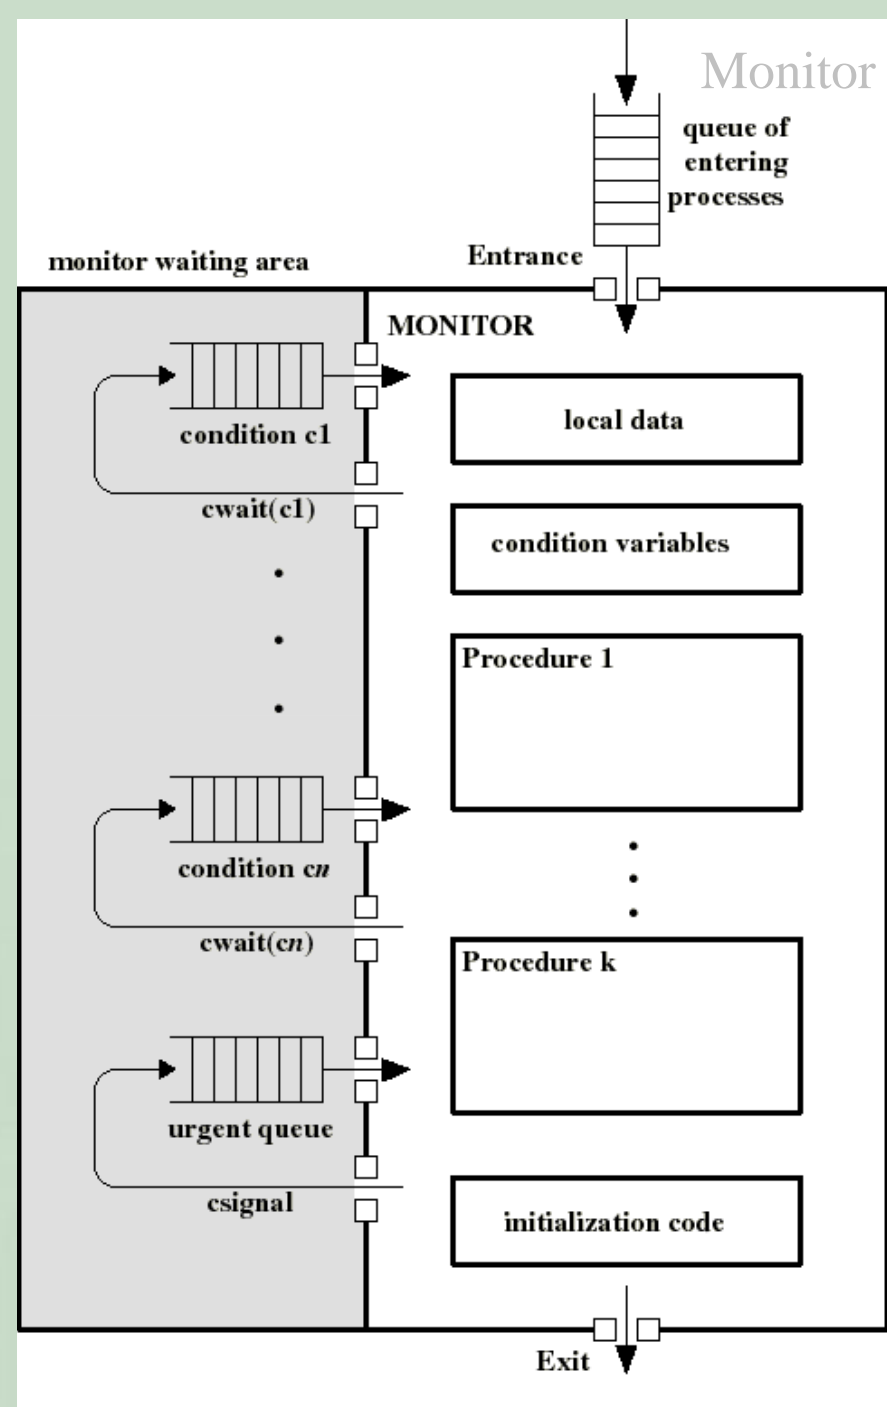
Monitor Mutual Exclusion

- The monitor ensures mutual exclusion - no need to program this constraint explicitly.
- The monitor **locks** (protects) shared data on process entry.
- Process synchronization is done by the programmer by using **condition variables**.
 - ⌘ conditions needing to be satisfied before entering monitor
 - ⌘ local to the monitor - accessible only within the monitor
 - **cwait(a)**: blocks execution of the calling process on condition variable **a**. The process can resume execution only if another process executes **csignal(a)**
 - **csignal(a)**: resume execution of some process blocked on condition variable **a**.
 - ⌘ If several such processes exist: choose any one
 - ⌘ If no such process exists: do nothing



Monitor

- Waiting processes are
 - in the entrance queue or
 - in a condition queue
- A process puts itself into condition queue c_i by issuing $cwait(c_i)$
- $csignal(c_i)$ brings into the monitor one process in condition c_i queue
- $csignal(c_i)$ blocks the calling process and puts it in the urgent queue
 - unless $csignal$ is the last operation of the monitor procedure



Monitor implementation

- Semaphore to enter monitor
- Semaphore for each condition variable
- Must allow multiple processes in the monitor
 - ☞ only one active
 - ☞ others are in condition variable queues



Monitor for the P/C problem

- Monitor holds the buffer:
 - ✧ buffer: array[0..k-1] of items;
- Two condition variables:
 - ✧ notfull: csignal(notfull) indicates that the buffer is not full
 - ✧ notempty: csignal(notempty) indicates that the buffer is not empty
- Buffer pointers and counts:
 - ✧ nextin: points to next item to be appended
 - ✧ nextout: points to next item to be taken
 - ✧ count: holds the number of items in buffer



Monitor for the P/C problem

Monitor boundedbuffer:

```
item buffer[0..k-1];  
int nextin=0, nextout=0, count=0;  
condition notfull, notempty;
```

```
Append(item v){  
    if (count==k) cwait(notfull);  
    buffer[nextin] = v;  
    nextin = (nextin + 1) mod k;  
    count++;  
    csignal(notempty);  
}
```

```
Take(){  
    if (count==0) cwait(notempty);  
    v = buffer[nextout];  
    nextout = (nextout + 1) mod k;  
    count--;  
    csignal(notfull);  
    return v;  
}
```



Monitor for the P/C problem

```
procedure producer() {  
    while (true) {  
        item v = produceItem();  
        boundedbuffer.Append(v);  
    }  
}  
procedure consumer() {  
    while (true) {  
        item v = boundedbuffer.Take();  
        consumeItem(v);  
    }  
}
```





Classical Synchronization Problems

Readers and Writers Problem

- Data object is shared (file, memory, registers)
 - ↻ many processes that only read data (readers)
 - ↻ many processes that only write data (writers)
- Conditions needing to be satisfied:
 - ↻ many can read at the same time (concurrency)
 - ↻ only one writer at a time and no one allowed to read while someone is writing (mutual exclusion)
- Solutions result in reader or writer priority



Readers/Writers (priority?)

```
Semaphore rmutex=1, wmutex = 1;  
integer readcount = 0;
```

```
WRITERS: while(true)  
{ wait(wmutex);  
  <write to the data object>  
  signal(wmutex);  
};
```

Only one writer
at a time

The first reader
makes sure no one
can write

```
READERS: while(true)  
{ wait(rmutex);  
  readcount++;  
  if (readcount == 1) wait(wmutex);  
  <read the data>  
  wait(rmutex);  
  readcount--;  
  if (readcount == 0) signal(wmutex);  
  signal(rmutex);  
};
```

More than one
reader at a time

Last one out allows
writing again

Writers/Readers (priority?)

Semaphore `outerQ = rsem = rmutex = wmutex = wsem = 1;`

```
while(true)
{ wait(outerQ);
  wait(rsem);
  wait(rmutex);
  readcnt++;
  if (readcnt == 1)
    wait(wsem);
  signal(rmutex);
  signal(rsem);
  signal(outerQ);
```

Additional readers
queue here allowing
writers to jump
ahead of the readers

Disable
writers

Once a writer wants to
write – no new readers
allowed

Last reader out
allows writers

READ

```
wait(rmutex);
readcnt--;
if(readcnt == 0)
  signal(wsem);
signal(rmutex);
};
```


```
while(true)
{ wait(wmutex);
  writecnt++;
  if (writecnt == 1)
    wait(rsem);
  signal(wmutex);
  wait(wsem);
  WRITE
  signal(wsem);
  wait(wmutex);
  writecnt--;
  if (writecnt == 0)
    signal(rsem);
  signal(wmutex);
};
```

Wait here until
all readers done

Last writer out
allows readers

Barbershop Problem

- 3 barbers, each with a barber chair
 - ✧ Haircuts vary in time
- Sofa can hold 4 customers
- Maximum of 20 customers in shop
 - ✧ Customers wait outside if necessary
- When a chair is empty:
 - ✧ Customer sitting longest on sofa is served
 - ✧ Customer standing the longest sits down on sofa
- After haircut, customer pays cashier at cash register
 - ✧ Algorithm has a separate cashier, but often barbers also take payment



Fair Barbershop

```

procedure customer;
var custnr: integer;
begin
  wait ( max_capacity );
  /* enter_shop */
  wait( mutex1 );
  count := count + 1;
  custnr := count;
  signal( mutex1 );
  wait( sofa );
  /* sit on sofa */
  wait( barber_chair );
  /* get up from sofa */
  signal( sofa );
  /* sit in barber chair */
  wait( mutex2 );
  enqueue( custnr );
  signal( cust_ready );
  signal( mutex2 );
  wait( finished[custnr] );
  /* leave barber chair */
  signal( leave_b_chair );
  /* pay */
  signal( payment );
  wait( receipt );
  /* exit shop */
  signal( max_capacity );
end;

```

```

procedure barber;
var b_cust: integer
begin
  repeat
    wait( cust_ready );
    wait( mutex2 );
    dequeue( b_cust );
    signal( mutex2 );
    wait( coord );
    /* cut hair */
    signal( coord );
    signal( finished[b_cust] );
    wait( leave_b_chair );
    signal( barber_chair );
  forever
end;

```

program
var

```

procedure cashier;
begin
  repeat
    wait( payment );
    wait( coord );
    /* accept payment */
    signal( coord );
    signal( receipt );
  forever
end;

```

```

barbershop;
max_capacity: semaphore (:=20);
sofa: semaphore (:=4);
barber_chair, coord: semaphore (:=3);
mutex1, mutex2: semaphore (:=1);
cust_ready, leave_b_chair, payment, receipt: semaphore (:=0);
finished: array [1..50] of semaphore (:=0);
count: integer;

```

