

Review

- Computer Overview
- OS Overview
- OS Structures
- Process and Process Control
- Thread
- Scheduling
- Concurrent Programming and Mutual Exclusion



Computer Systems

- Registers
- Interrupts
- Caching
- Memory Hierarchy
- Input/output
- Protection



Operating Systems

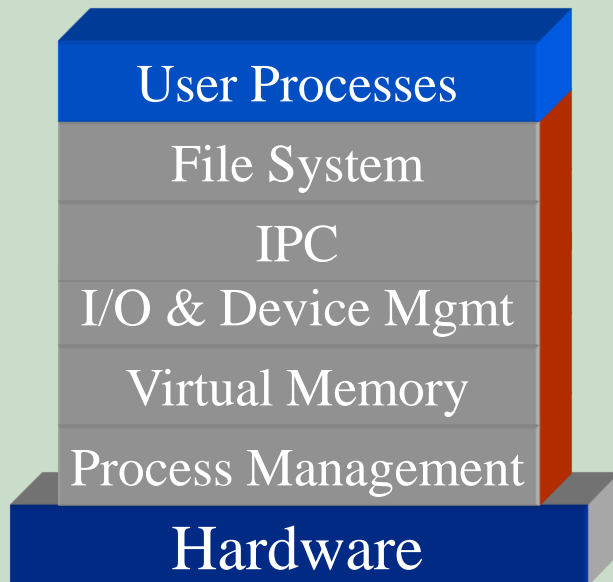

- Objectives: Convenience, Efficiency and Ability to evolve.
 - ❧ Process and Thread
 - ❧ Scheduling
 - ❧ Concurrency and Mutual exclusion
 - ❧ Resource management
 - ❧ Memory Management
 - ❧ File System
 - ❧ Communication




OS Structures

- Layered Structure
- Multi-kernels for Multi-core processors
- Microkernel Structure (IPC)
- Virtual Machine Structure





Traditional OS



Microkernel OS

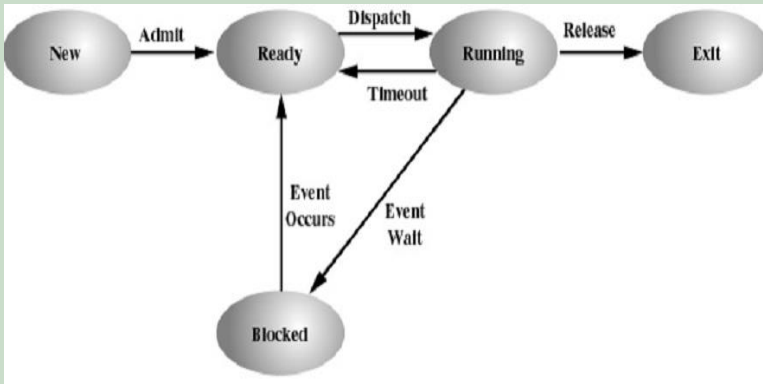


Process Concepts

- Process Definition
- Process Creation
- Process States and State Transitions
- Important Information Associated with Processes (PCB)




State Transitions



+ Suspend State

Five States



Process Control

- Process Control Block (PCB)
- Process Execution
- Context Switch
- Process Scheduling



When to Switch a Process

- Clock interrupt
- I/O interrupt
- Memory fault
- Trap (error occurred)
- Supervisor call (system calls)
- Context Switch:
 - ∞ Saving the state of one process and loading another process onto the CPU

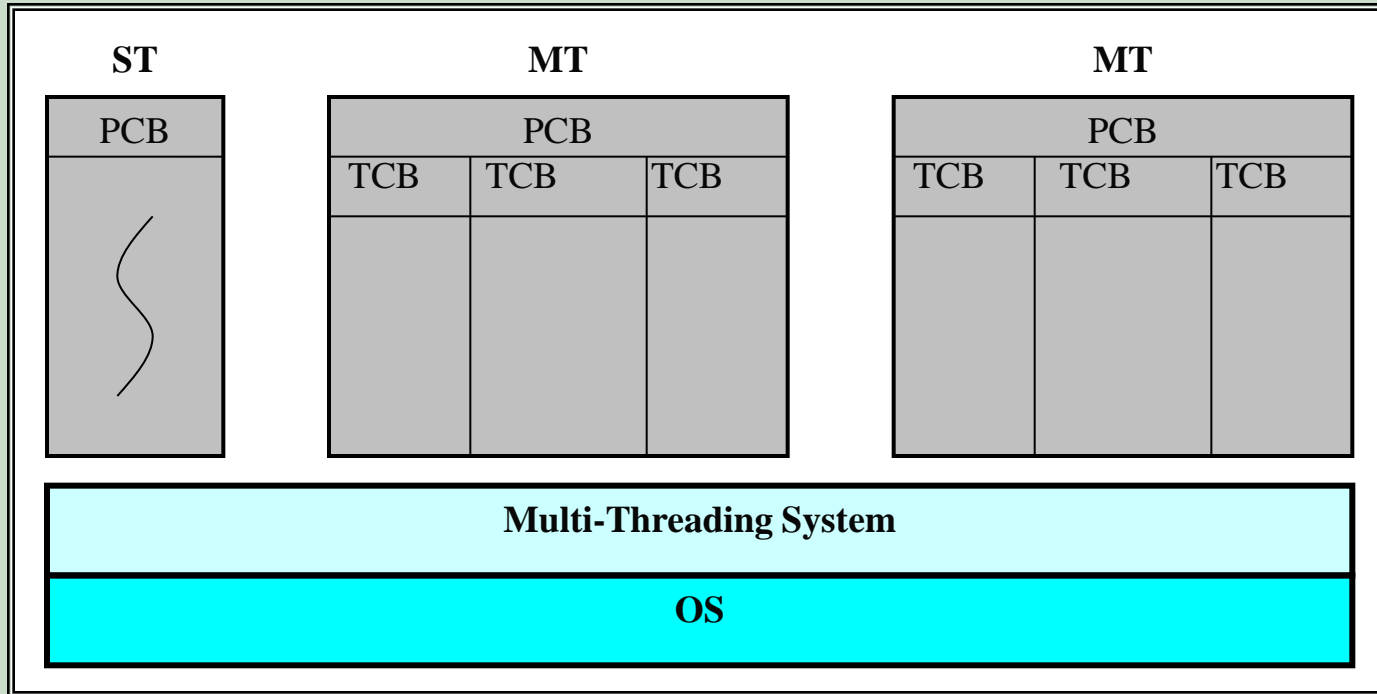


Threads

- A thread consists of:
 - ⌘ state (Running, Ready, etc.)
 - ⌘ CPU context (program counter, register set.)
 - ⌘ an execution stack.
 - ⌘ access to the memory and resources of its process (shared with all other threads in that process.)
 - ⌘ OS resources (open files, signals, etc.)
- Benefits:
 - ⌘ Far less time to create/terminate.
 - ⌘ Switching between threads is faster.
 - ⌘ No memory management issues, etc.



Threads and Processes



PCB : Process Control Block

TCB : Thread Control Block

contains information and status of thread/process

Scheduling

- Objectives:
 - ↻ minimize response time
 - ↻ maximize throughput
 - ↻ maximize processor efficiency
- Long-term
 - ↻ performed when new process is created
- Medium-term
 - ↻ swapping to maintain a degree of multiprogramming
- Short-term
 - ↻ which ready process to execute next – dispatcher



Preemptive vs. Non-preemptive

- Scheduling that only takes place due to I/O or process termination is non-preemptive.
- Preemptive scheduling allows the operating system to interrupt the currently running process and move it to the ready state.
- Preemptive scheduling:
 - ✧ incurs greater overhead (context switch)
 - ✧ provides better service to the total population of processes
 - ✧ may prevent one process from monopolizing the processor



Algorithms

- First Come First Served
 - ∞ Processes queued in order of arrival
 - ∞ Runs until finished or blocks on I/O
- Round Robin
 - ∞ FCFS with preemption
 - ∞ Size of ticks affects performance
- Shortest Process Next
 - ∞ Select process with shortest expected running time (non-preemptive)
 - ∞ Difficult to estimate required time



Algorithms

- Highest Response Ratio Next
 - ✧ Non-preemptive, tries to get best average normalized turnaround time
 - ✧ Depends on Response Ratio
 - W = time spent waiting
 - S = expected service time $RR = (W + S) / S$
- Priority
 - ✧ Schedule Process with the highest priority
- Feedback Queue
 - ✧ Starts in high-priority queue, moves down in priority as it executes
 - ✧ Lower-priority queues often given longer time slices



FCFS

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- Suppose that the processes arrive at time 0 in the order:
 P_1, P_2, P_3

The simplified Gantt Chart for the schedule is:




- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: $(0 + 24 + 27)/3 = 17$



RR with Time Quantum = 20

<u>Process</u>	<u>Burst Time</u>
P_1	53
P_2	17
P_3	68
P_4	24

- The Gantt chart is:



Resource Competition

■ Mutual Exclusion


- ⌘ Critical resource – a single non-sharable resource.
- ⌘ Critical section – portion of the program that accesses a critical resource.

■ Deadlock

- ⌘ Each process owns a resource that the other is waiting for.
- ⌘ Two processes are waiting for communication from the other.

■ Starvation

- ⌘ A process is denied access to a resource, even though there is no deadlock situation.



Mutual Exclusion and Synchronization

- Software Solutions:
 - ☞ Shared variables, take turns
- Hardware Solutions:
 - ☞ Special Instructions: test&set, swap
- OS Solutions:
 - ☞ Semaphores: Binary or Counting
 - ☞ Operations: P/V, wait/signal, lock/unlock
 - ☞ Monitors



A simple synchronization example

Expression: $(x^2) / (y-z) + \sin(x)$

Shared variables: t_1, t_2, t_3, t_4, t_5 ;


P1: $t_1 = x^2$;

P2: $t_2 = y - z$;

P3: $t_3 = \sin(X)$;

P4: $t_4 = t_1/t_2$;

P5: $t_5 = t_4 + t_3$



Concurrency control

semaphore s1, s2; // initial 0s


P1: $t1 = x^2$; $V(s1)$;

P2: $t2 = y - z$; $V(s1)$;

P3: $t3 = \sin(X)$; $V(s2)$;

P4: $P(s1)$; $P(s1)$; $t4 = t1/t2$; $V(s2)$;

P5: $P(s2)$; $P(s2)$; $t5 = t4 + t3$;



Bounded Buffer Solution

Shared semaphore: `empty = n, full = 0, mutex = 1;`
 Item buffer[n], `int in = out = 0;`

repeat
 produce an item in nextp

`wait(empty);`
`wait(mutex);`

`buffer[in] = nextp`
`in = (in + 1) mod n`

`signal(mutex);`
`signal(full);`

until false

Producer

repeat

`wait(full);`
`wait(mutex);`

`nextc = buffer[out]`
`out = (out + 1) mod n`

`signal(mutex);`
`signal(empty);`

consume the item in nextc

until false

Consumer

